


# STORAGE HIERARCHY



## Snowflake Accounts

- A customer can have as many accounts as they want
- Each account has its own URL – which uses the name of the account.
- Each account is deployed on a single cloud provider platform (AWS, Azure, GCP)
- Each account exists in a single geographic region
- Each account exists with a single Snowflake Edition (e.g. Standard, Premier, etc.)

## Accounts Contain Databases

- Each database belongs to a single Snowflake account
- Databases can be replicated to other accounts, but they cannot SPAN multiple accounts


## Databases Contain Schemas

- Each schema belongs to a single Snowflake database, in a single account
- Schemas can be replicated to other accounts or databases, but they cannot SPAN accounts or databases

## Schemas Contain Other Objects

- Objects include tables, views, file formats, sequences, UDFs, & stored procs
- Objects belong to a single schema, in a single database, in a single account

# STORAGE CONTAINER HIERARCHY



## Snowflake Accounts

The account shown here is the account NS61174. Most trial accounts start with two letters followed by 5 numbers. Trial accounts can be converted to paid accounts.

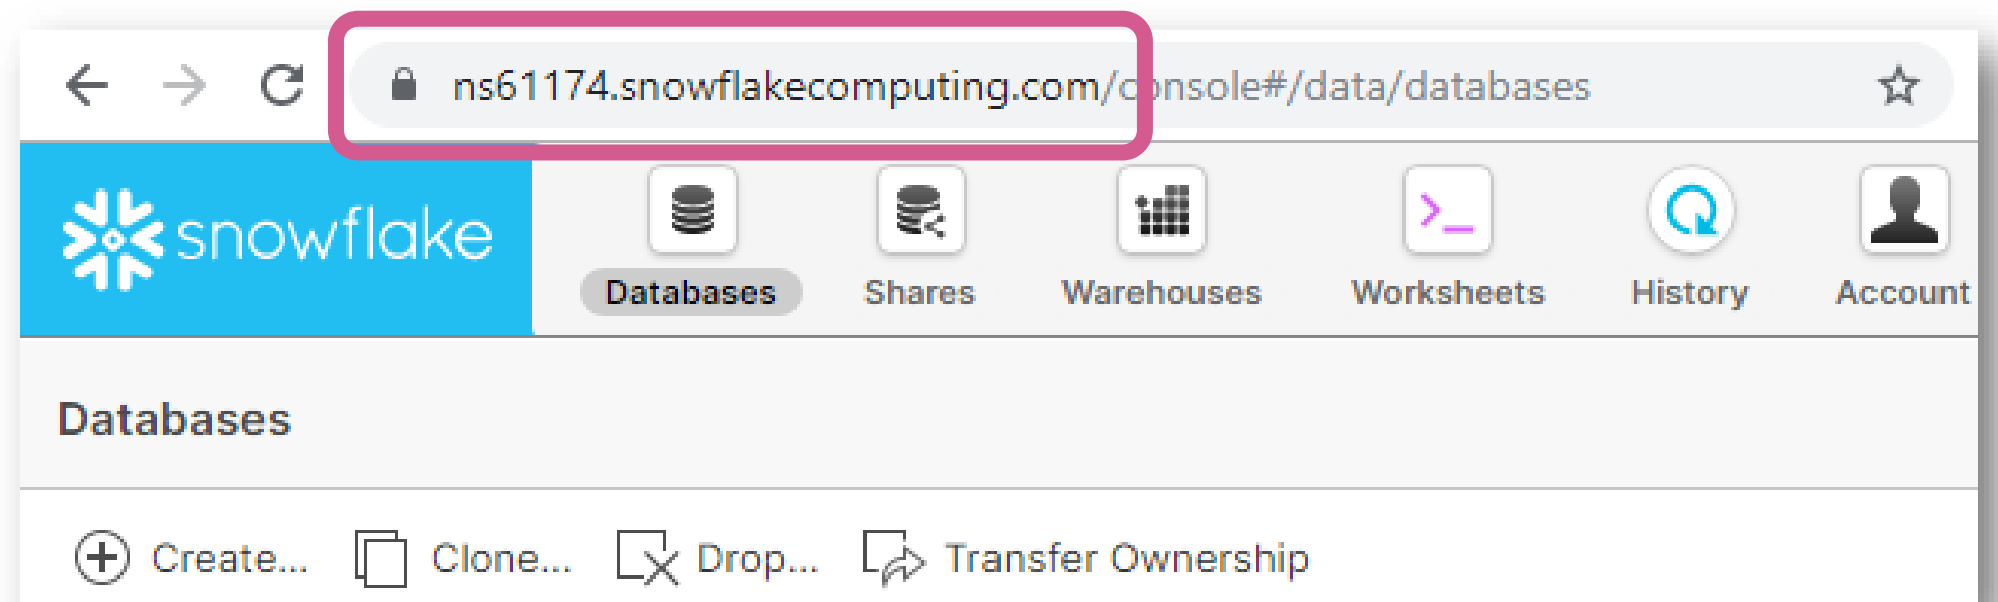
Customers can request a “vanity” address like:

<https://robertsfishandchips.snowflakecomputing.com>

When an account has just the account name, followed by “snowflakecomputing.com”, we know the cloud provider is **AWS** and the region is **us-west-2**. This is because this was the first cloud provider and region on which Snowflake deployed.

Later, when Snowflake expanded to other **regions** and **providers**, the region name was added to the account URL. For example, an address like:


<https://xy12345.east-us-2.azure.snowflakecomputing.com>



READ MORE ABOUT REGIONS HERE:

<https://docs.snowflake.net/manuals/user-guide/intro-regions.html>

# STORAGE CONTAINER HIERARCHY




## Accounts Contain Databases

- In this example, account NS61174 has three databases.
- The **SNOWFLAKE** database contains information about account usage. It is automatically added by Snowflake to each new account. It is sometimes called the “Account Usage Share” because it is shared by Snowflake so customers can check usage.
- The **SNOWFLAKE\_SAMPLE\_DATA** is also shared by Snowflake. This database allows for performance testing and tutorials. Customers are not charged for storing this data., because they are NOT storing it. The database is a shared database.

The screenshot shows the Snowflake web console interface. The browser address bar is highlighted with a pink box, showing the URL `ns61174.snowflakecomputing.com/console#/data/databases`. The console header includes the Snowflake logo and navigation tabs for Databases, Shares, Warehouses, Worksheets, History, and Account. The 'Databases' tab is active. Below the header, there are action buttons: '+ Create...', 'Clone...', 'Drop...', and 'Transfer Ownership'. A table lists the databases in the account. The table has columns for Database, Origin, Creation Time, and Owner. The first row, 'MY\_DATABASE', is highlighted with a purple box. The second row, 'SNOWFLAKE\_SAMPLE\_DATA', and the third row, 'SNOWFLAKE', are highlighted with a light blue background.

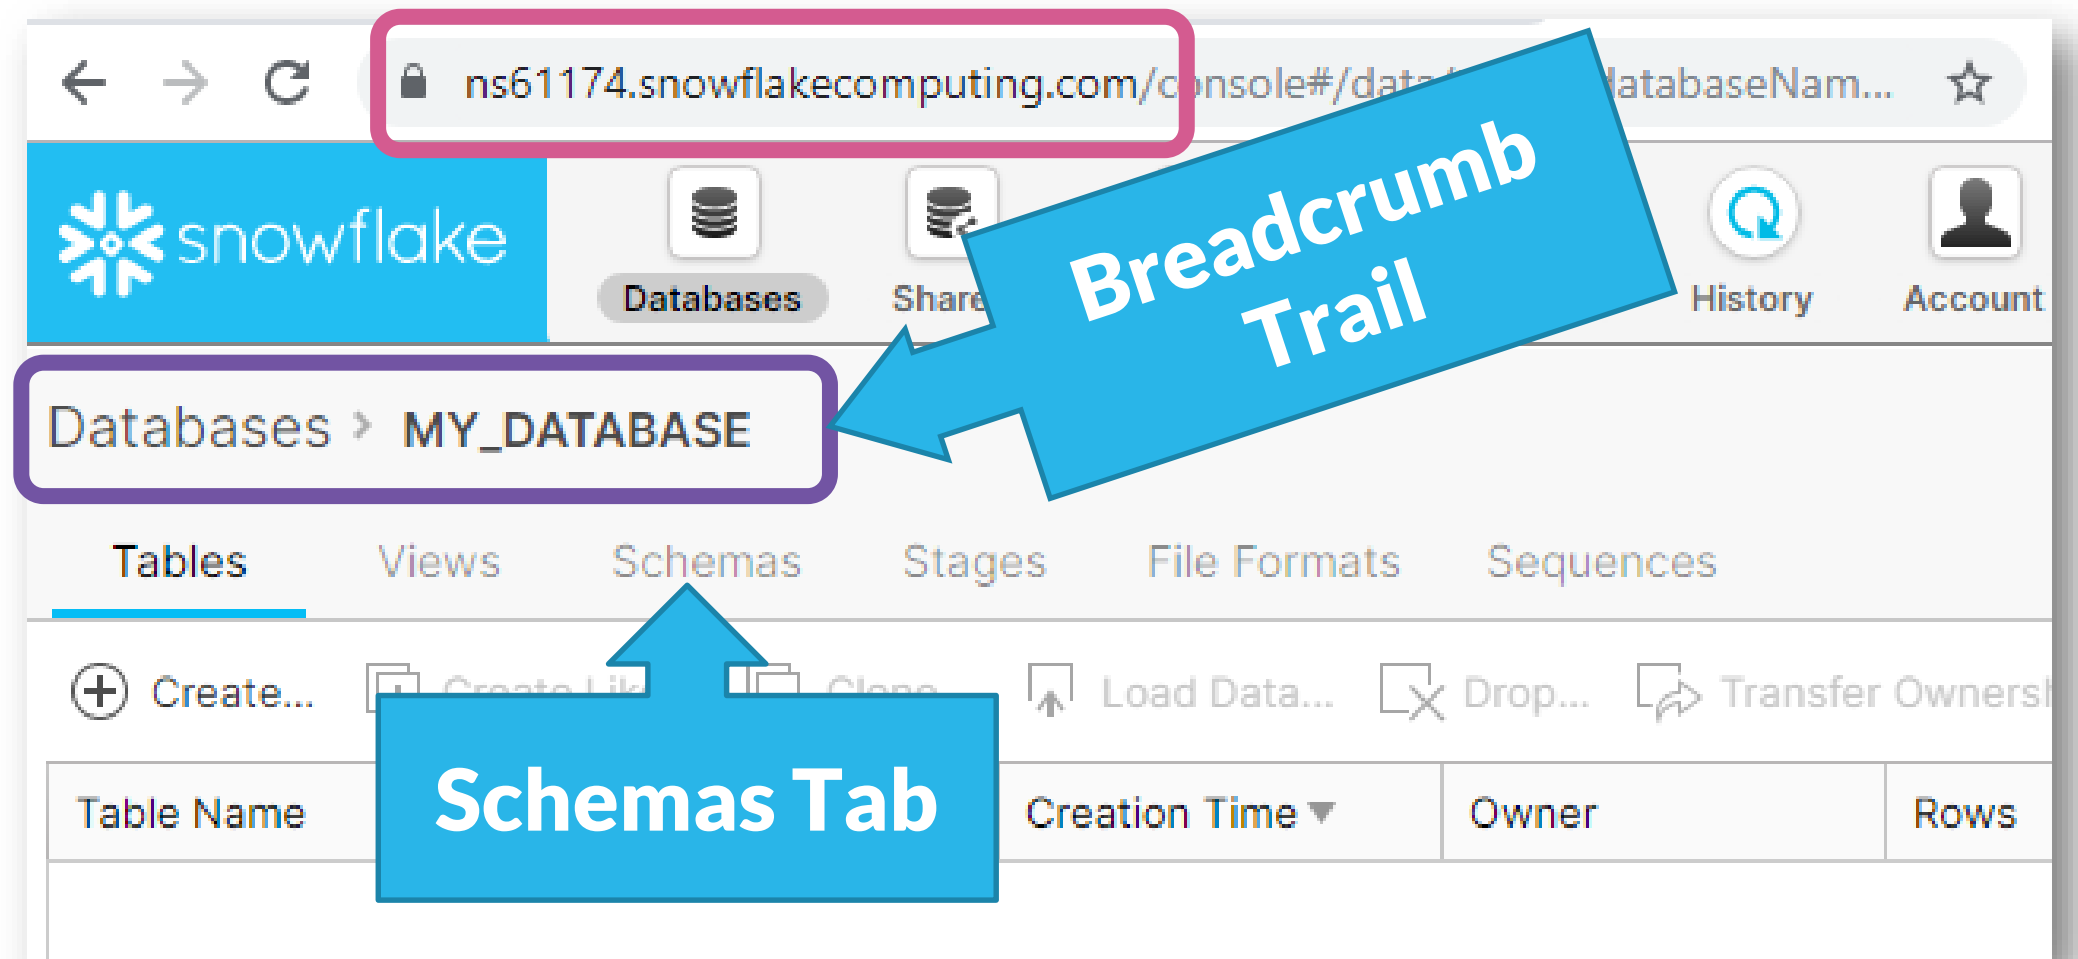
Database	Origin	Creation Time ▼	Owner
MY_DATABASE		7:06:25 AM	ACCOUNTADMIN
SNOWFLAKE_SAMPLE_DATA	SFC_SAMPLES.SA...	6/1/19 10:51:59 PM	ACCOUNTADMIN
SNOWFLAKE	SNOWFLAKE.ACCO...	6/1/19 10:51:47 PM	

# STORAGE CONTAINER HIERARCHY




## Accounts Contain Databases

When we drill into any database, the name of the database gets added to the breadcrumb trail and a set of tabs appears showing Schemas and other object types.



# STORAGE CONTAINER HIERARCHY



Database.Schema = Namespace

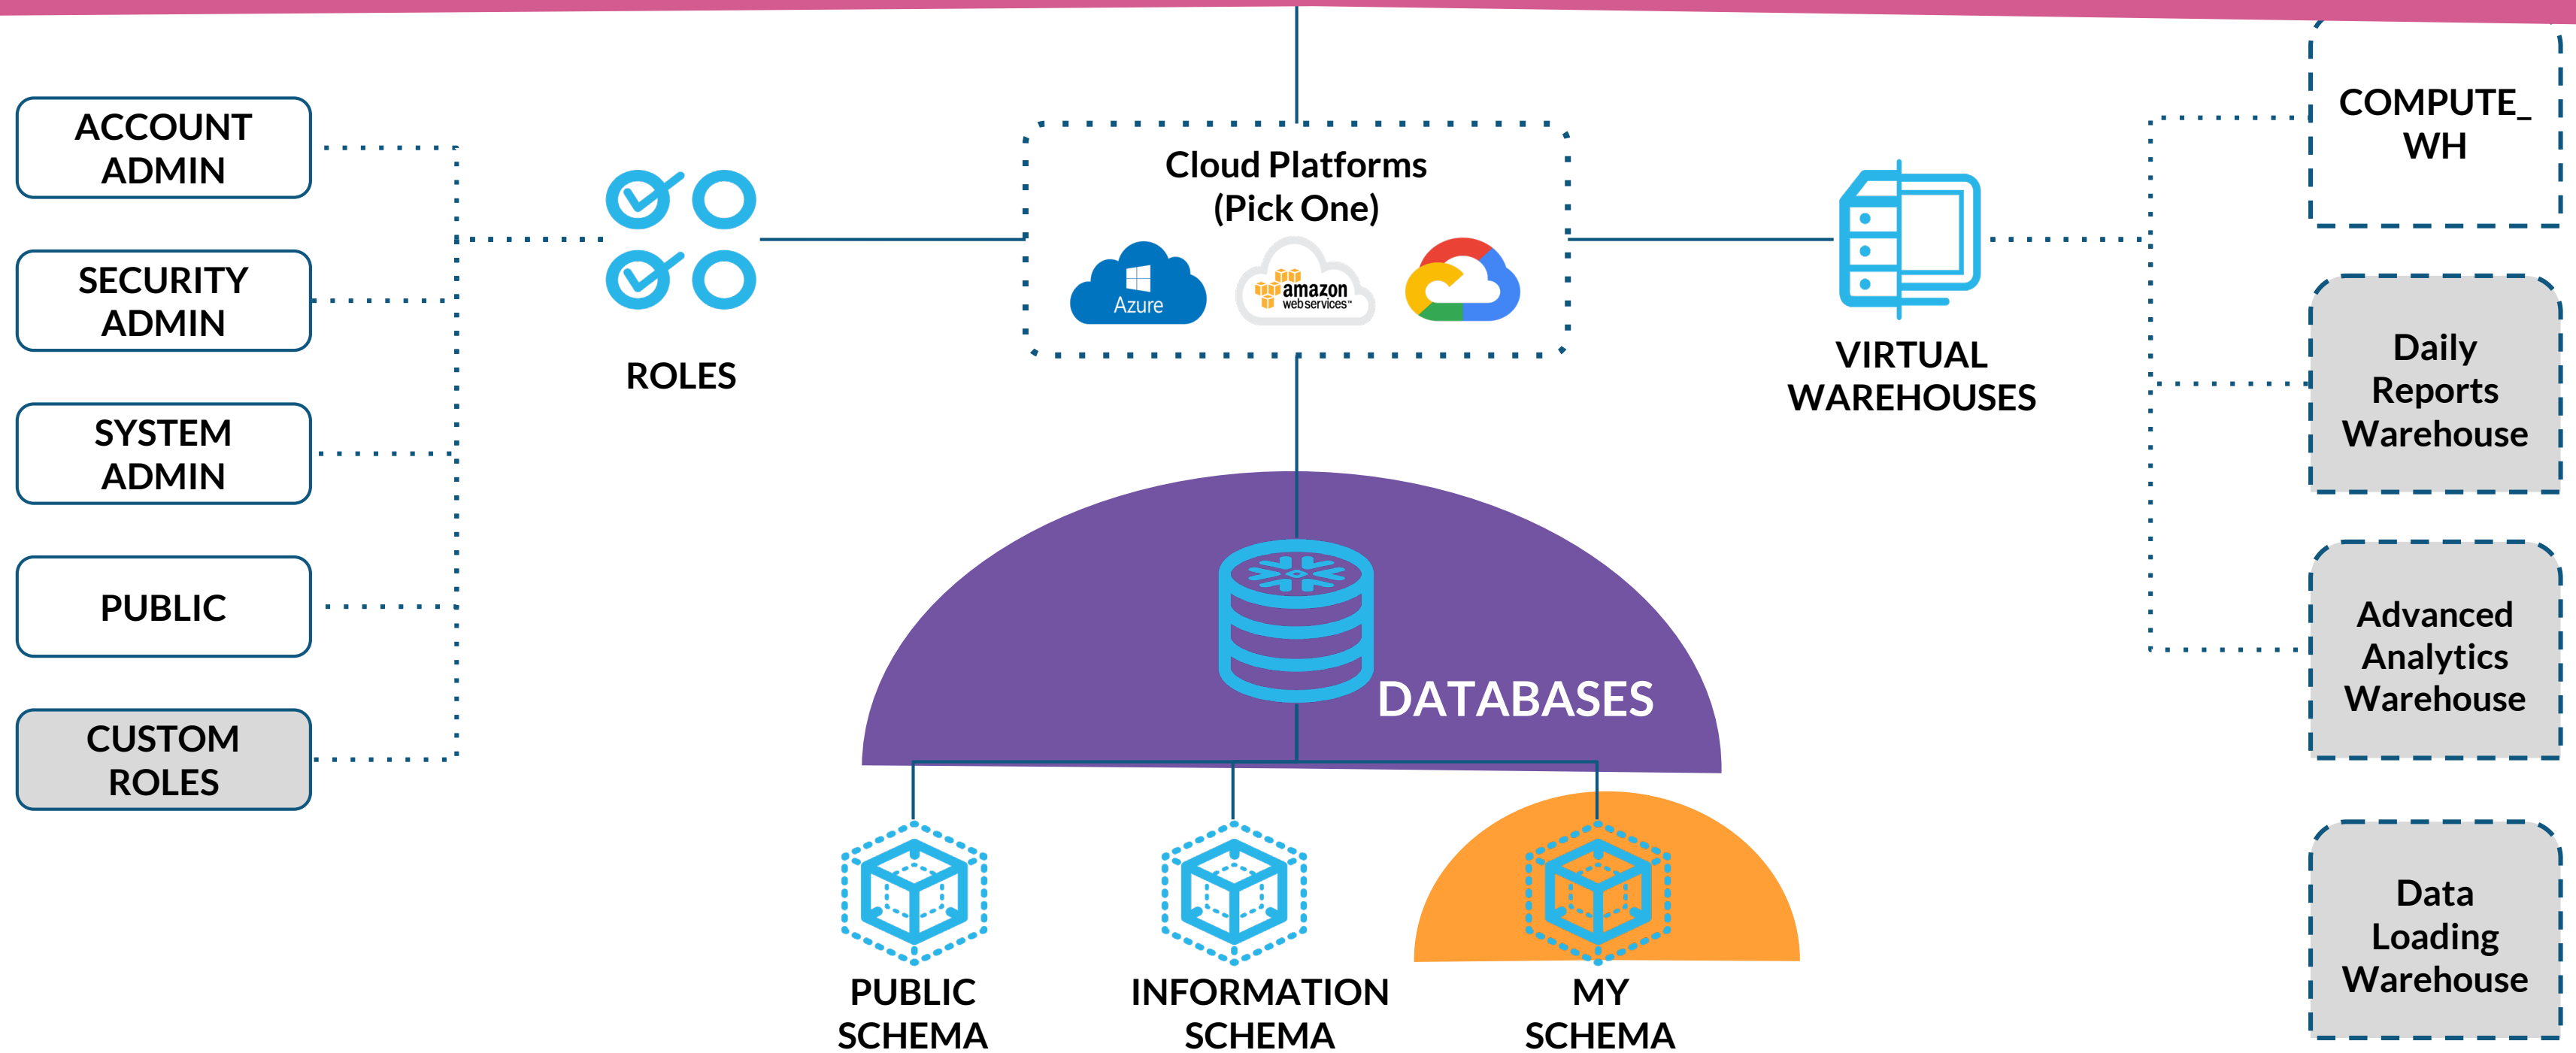
## Databases Contain Schemas

- The set of tabs below the breadcrumb trail makes it look as though Schemas are at the same level as Tables, Views, Stages, File Formats, and Sequence but this is not an accurate way to think of Schemas.
- Schemas are a logical grouping one level above the other objects listed.

A screenshot of the Snowflake console interface. The browser address bar shows the URL `ns61174.snowflakecomputing.com/console#/data/schemas?databaseN...`. The console header includes the Snowflake logo and navigation tabs: Databases, Shares, Warehouses, Worksheets, History, and Account. The breadcrumb trail shows 'Databases > MY\_DATABASE'. Below the breadcrumb, there are tabs for 'Tables', 'Views', 'Schemas' (which is selected), 'Stages', 'File Formats', and 'Sequences'. A toolbar contains actions: '+ Create...', 'Clone...', 'Alter...', 'Drop...', and 'Transfer Ownership'. A table lists the schemas in the database:


Schema	Creation Time ▼	Owner	Managed Access	Comment
INFORMATION_SCHEMA	7:21:55 AM			
MY_SCHEMA	7:06:39 AM	ACCOUNTADMIN		
PUBLIC	7:06:25 AM	ACCOUNTADMIN		

# SNOWFLAKE ACCOUNT





# STORAGE CONTAINER HIERARCHY



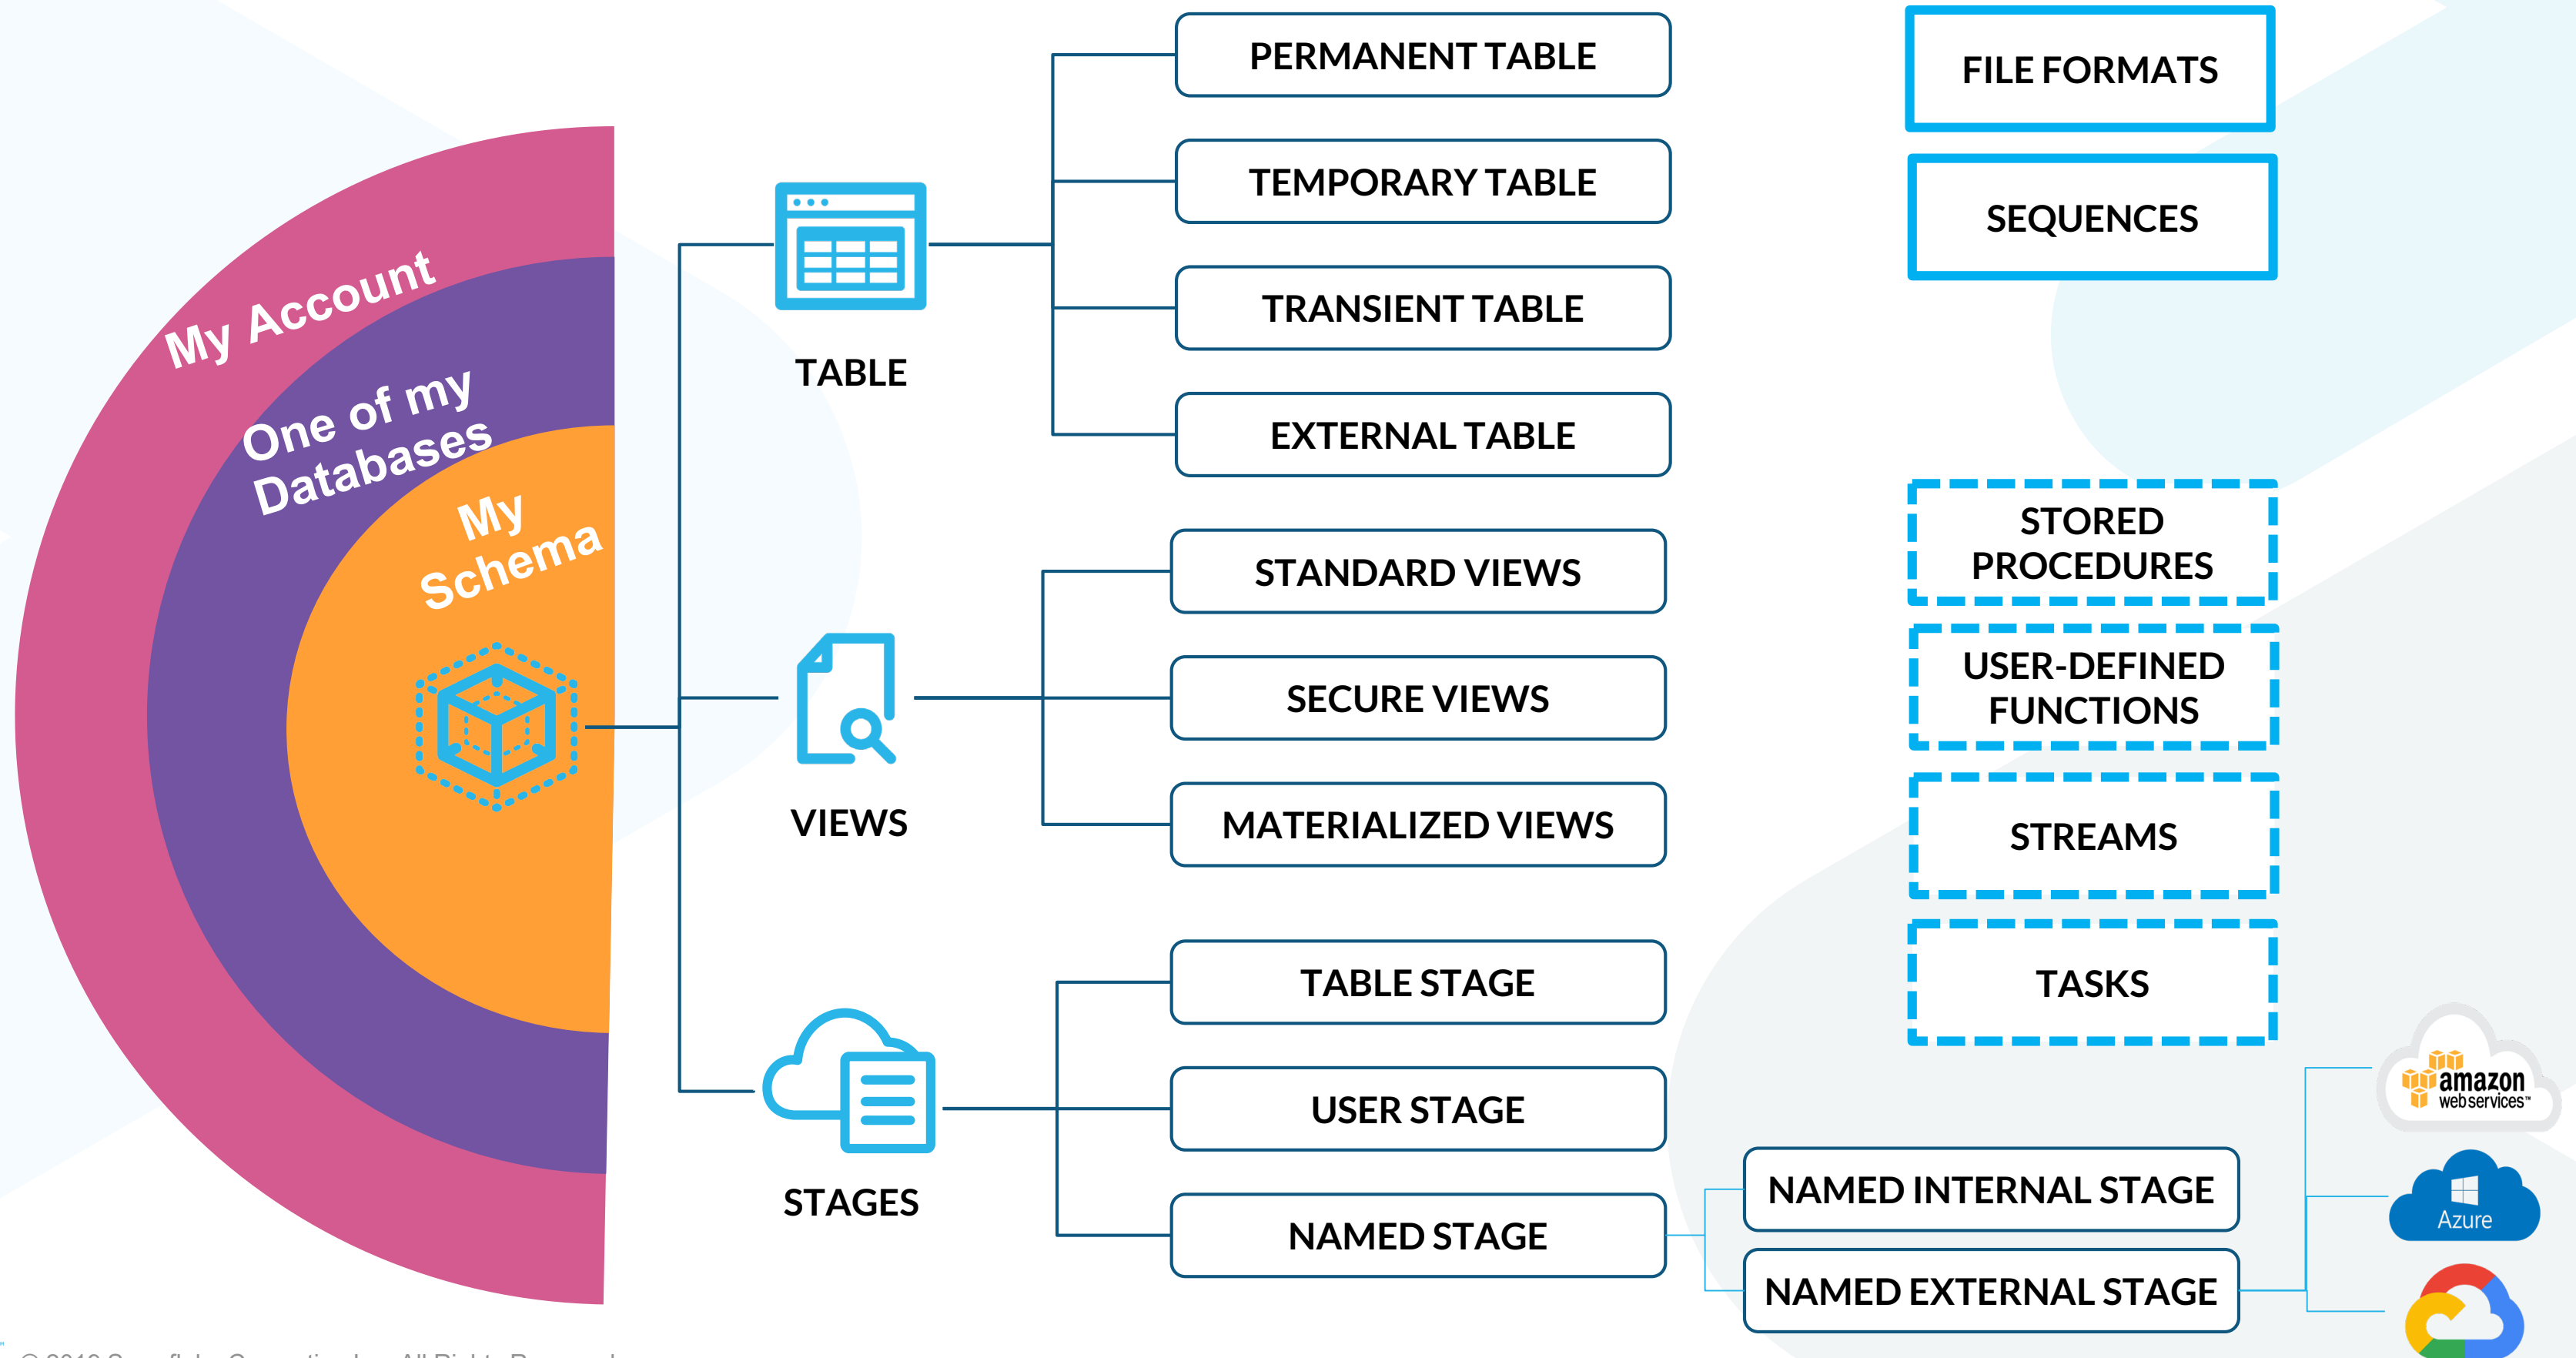
## Tables, Views, Stages, File Formats, and Sequences are situated within Schemas

- The tables shown below are in different schemas.
- Snowflake allows you to bypass the schema during the drill-down process, making it seem as if tables and other objects are not within a schema, but they are.

The screenshot shows the Snowflake web interface. The top navigation bar includes the Snowflake logo and tabs for 'Databases', 'Shares', 'Warehouses', and 'Workspaces'. The 'Databases' tab is selected, and the breadcrumb 'Databases > MY\_DATABASE' is shown. Below the breadcrumb, there are tabs for 'Tables', 'Views', 'Schemas', 'Stages', 'File Formats', and 'Sequences'. The 'Tables' tab is selected. A toolbar contains actions: '+ Create...', '+ Create Like...', 'Clone...', 'Load Data...', and 'Drop'. Below the toolbar is a table listing tables in the 'MY\_DATABASE' database. The 'Schema' column is highlighted with an orange box, showing 'MY\_SCHEMA' and 'PUBLIC' for the two tables listed.

Table Name	Schema	Creation Time ▼
MY_SECOND_TABLE	MY_SCHEMA	9:00:11 AM
MY_FIRST_TABLE	PUBLIC	8:59:55 AM

# SCHEMA OBJECTS





# TABLE TYPES



## PERMANENT

- Standard, default table type.
- Most commonly used.
- Default Time-Travel
- 7 Day Fail-Safe – non-configurable



## TEMPORARY

- Lasts single session
- Can be used for moving data from one permanent to another with some transformation within the table
- No costs associated with failsafe storage



## TRANSIENT

- No costs associated with failsafe storage
- Usually less costly than permanent tables.
- Time travel up to one day is available



## EXTERNAL

- Stored in External Stages
- Read Only
- DML SQL works: Selects, Aggregations, Joins
- Can use in a View
- No cloning, time travel or fail-safe

Time-Travel



X


Fail-Safe



X




# VIEW TYPES




## Standard/Regular View

- Default View type
- Most commonly used
- Underlying DDL available to any role with access to the view



## Secure View


- The underlying query definition and details are only visible to authorized users
- Snowflake query optimizer behaves differently when running the SQL contained in a secure view so that it doesn't accidentally reveal things about the underlying data
- May not run as fast as a standard view, because the optimizer isn't used the same way



## Materialized View

- It has “view” in the name, but it behaves more like a table
- Results are stored just as a table would be
- In Snowflake, these are auto-refreshed
- Can provide cost savings but are inflexible and somewhat limited in use

# STAGE TYPES



**Table Stage**  
@%[TABLE\_NAME]

- Automatically defined – No set up needed
- Not viewable on Stages tab of WebUI
- You must specify certain file format-like parameters as part of the command during each COPY INTO event



**User Stage**  
@~[LOGIN]



**Named Internal Stage**  
@[STAGE\_NAME]

- Definable via the WebUI
- More flexible than Table and User Stages
- Viewable on Stages tab of Database Area of WebUI
- Have embedded default File Format definitions or you can specify named File Formats for use



**Named External Stage**  
@[STAGE\_NAME]

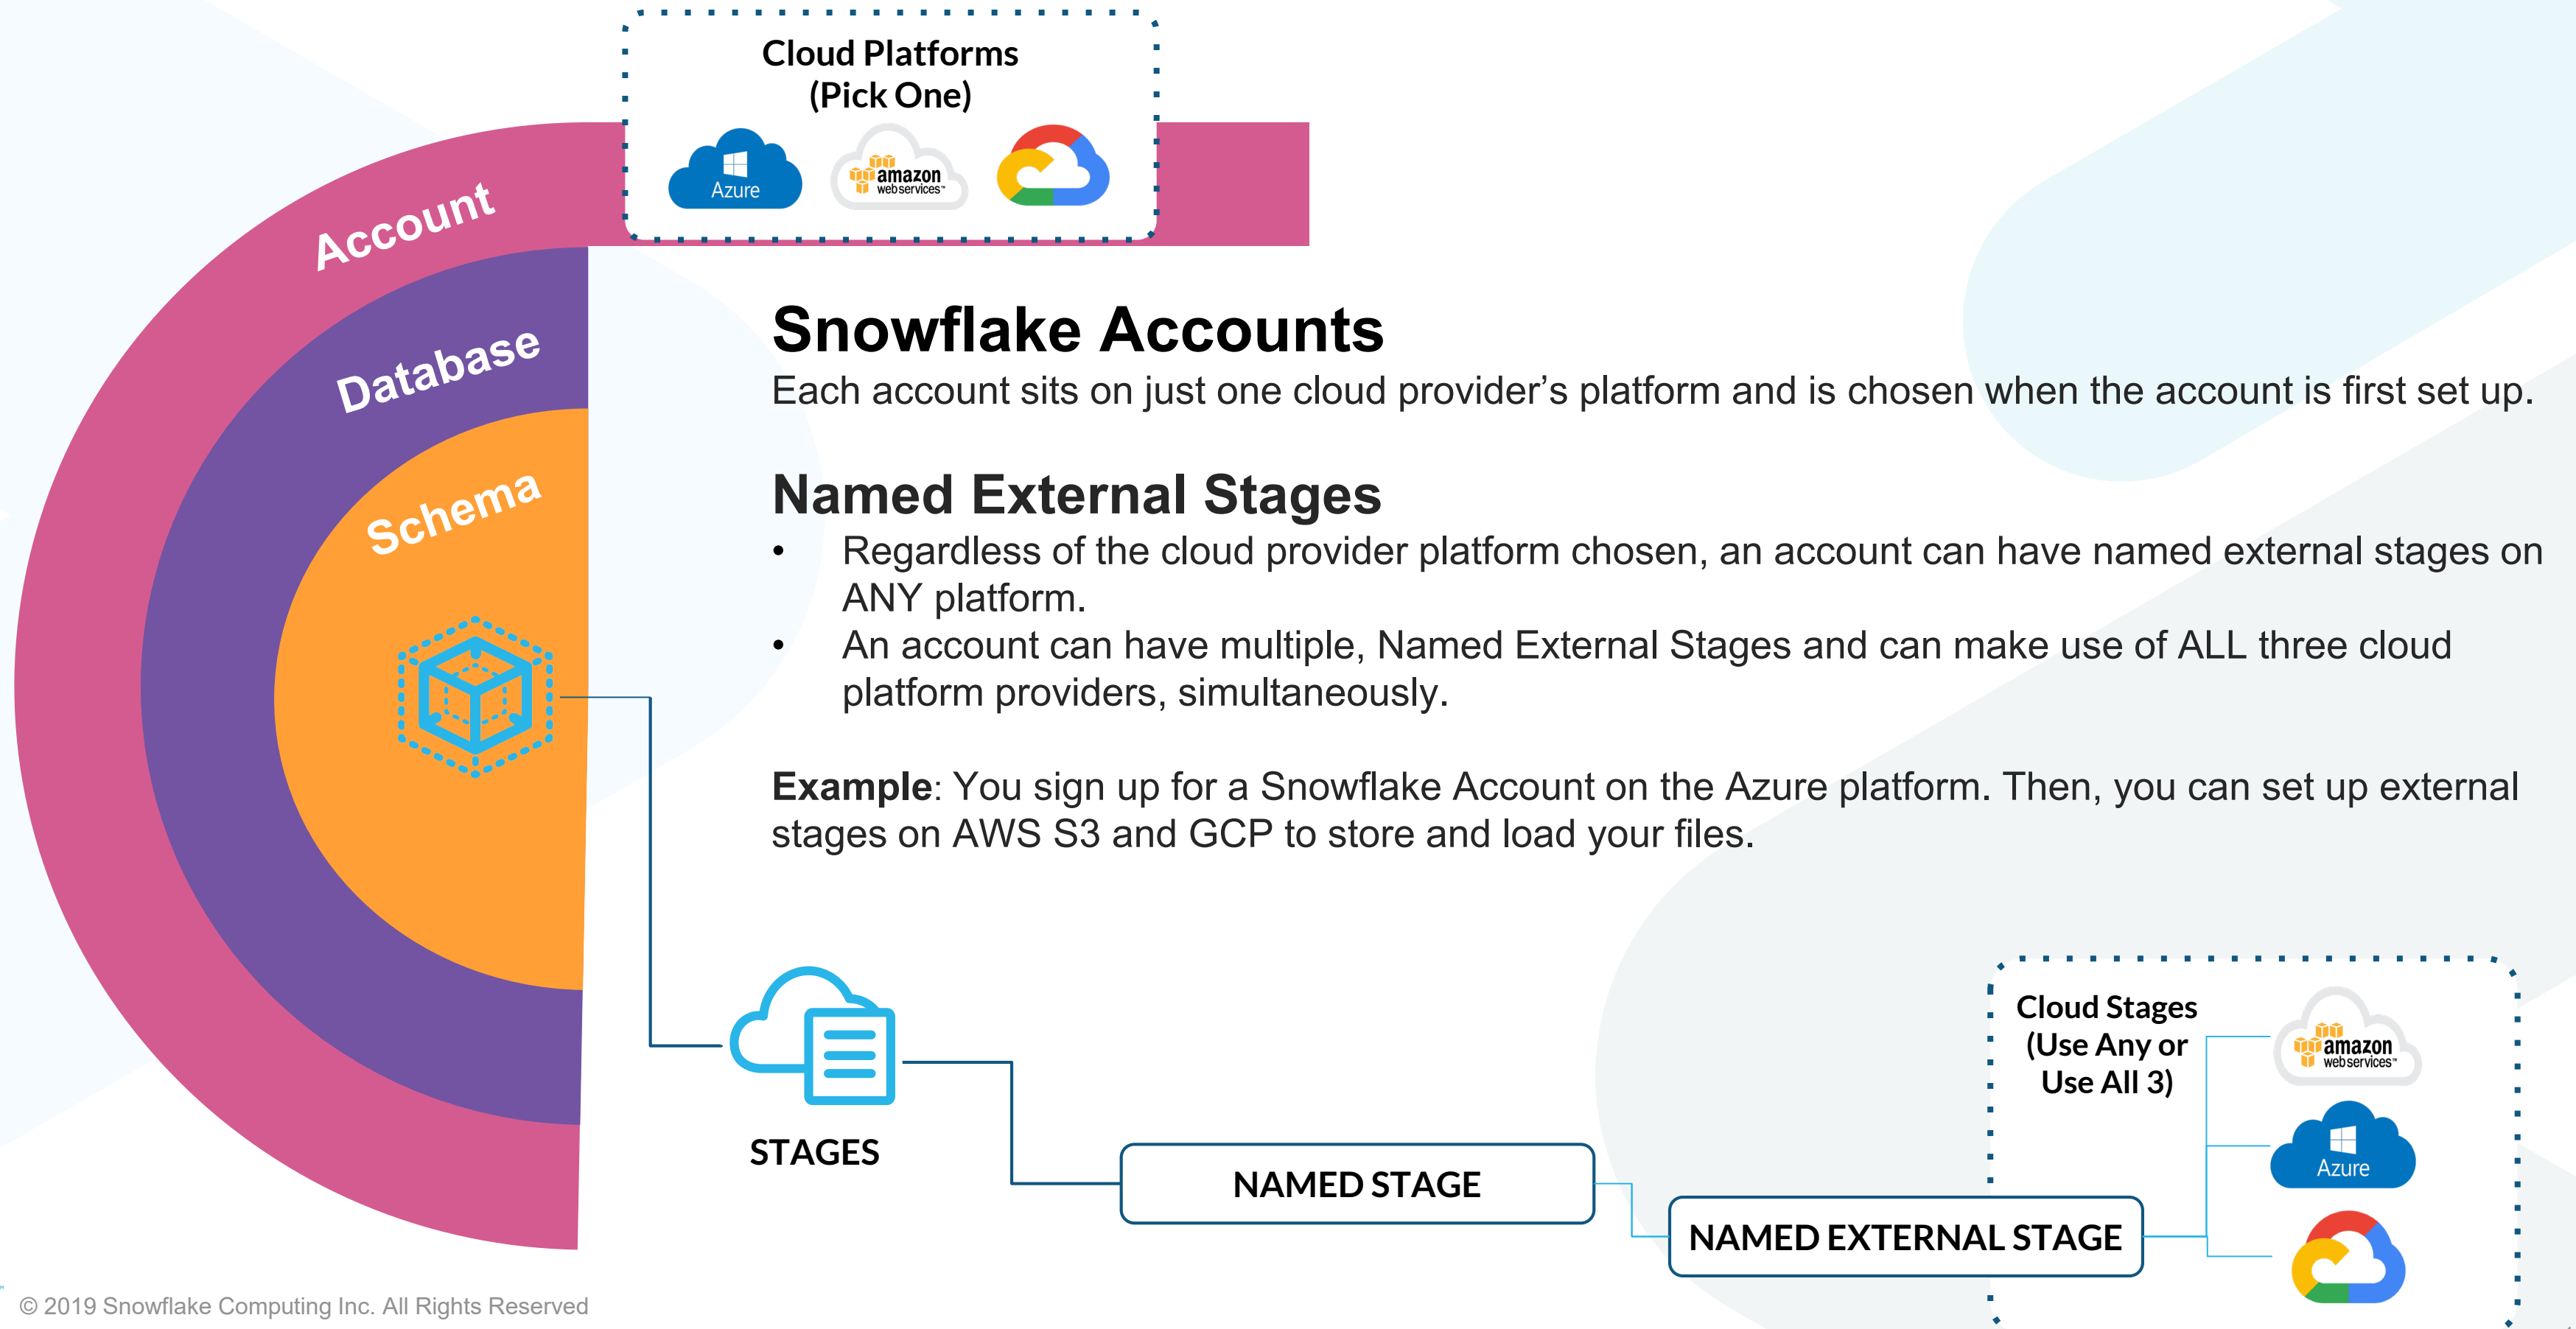
## GETTING FILES INTO AND OUT OF STAGES

- Use the PUT command to load files into any type of stage\*
- Use the COPY INTO command to move files from any stage into a table
- Use the COPY INTO command to unload files from tables and put them into any kind of stage.


\* The PUT command cannot be run from WebUI Worksheet. The PUT command can be run from SnowSQL, the Snowflake CLI client. The PUT command can also be run via other programmatic interfaces. You can also load files into a stage by going to the AWS, GCP, or Azure console and dragging them into storage or uploading them via the tools offered by the cloud platform providers.



# PLATFORM VS. STAGES



# STORAGE HIERARCHY



## Further Reading:

### DATABASES AND SCHEMAS

<https://docs.snowflake.net/manuals/sql-reference/ddl-database.html#database-schema-share-ddl>

### TABLES

<https://docs.snowflake.net/manuals/user-guide/tables-temp-transient.html#comparison-of-table-types>

### VIEWS

<https://docs.snowflake.net/manuals/user-guide/views-introduction.html>

### STAGES

<https://docs.snowflake.net/manuals/user-guide/data-load-local-file-system-create-stage.html#types-of-stages>

